

CIS 657 – Principles of Operating Systems

Topic: Introduction to Operating Systems

Endadul Hoque

Acknowledgement

- Youjip Won (Hanyang University)
- OSTEP book – by Remzi and Andrea Arpaci-Dusseau (University of Wisconsin)

What happens when a program runs?

- A running program executes instructions.
 1. The processor **fetches** an instruction from memory.
 2. **Decode**: Figure out which instruction this is
 3. **Execute**: i.e., add two numbers, access memory, check a condition, jump to function, and so forth.
 4. The processor moves on to the **next instruction** and so on.

Von Neumann model of computing

An Operating System (OS)

- Responsible for
 - Making it easy to **run** programs
 - Allowing programs to **share** memory
 - Enabling programs to **interact** with devices

OS is in charge of making sure the system operates **correctly** and **efficiently**.

How does OS ensure this?

VIRTUALIZATION

Virtualization

- The OS takes a physical resource and transforms it into a virtual form of itself.
 - **Physical resource:** Processor, Memory, Disk ...
- The virtual form is more general, powerful and easy-to-use.
- Sometimes, we refer to the OS as a **virtual machine**.

System Calls

- System call allows users **to tell the OS what to do**.
 - The OS provides some interface (APIs, standard library).
 - A typical OS exports a few hundred system calls.
 - Run programs
 - Access memory
 - Access devices

The OS as a resource manager

- The OS **manage resources** such as *CPU*, *memory* and *disk*.
- The OS allows
 - Many programs to run → Sharing the CPU
 - Many programs to *concurrently* access their own instructions and data → Sharing memory
 - Many programs to access devices → Sharing disks

Virtualizing the CPU

```
1      #include <stdio.h>
2      #include <stdlib.h>
3      #include <sys/time.h>
4      #include <assert.h>
5      #include "common.h"
6
7      int
8      main(int argc, char *argv[])
9      {
10         if (argc != 2) {
11             fprintf(stderr, "usage: cpu <string>\n");
12             exit(1);
13         }
14         char *str = argv[1];
15         while (1) {
16             Spin(1); // Repeatedly checks the time and
17                     // returns once it has run for a second
18             printf("%s\n", str);
19         }
20         return 0;
21     }
```

Simple Example(cpu.c): Code That Loops and Prints

Virtualizing the CPU

- Execution result 1

```
prompt> gcc -o cpu cpu.c -Wall
prompt> ./cpu "A"
A
A
A
^C
prompt>
```

Run forever; Only by pressing "Control-c" we can halt the program

Virtualizing the CPU

- Execution result 2

```
prompt> ./cpu A & ; ./cpu B & ; ./cpu C & ; ./cpu D &  
[1] 7353  
[2] 7354  
[3] 7355  
[4] 7356  
A  
B  
D  
C  
A  
B  
D  
C  
A  
...
```

Even though we have only
one processor, all four of
programs seem to be
running **at the same time!**

Virtualizing the CPU

- The system has a very large number of virtual CPUs.
 - Turning a single CPU into a seemingly infinite number of CPUs.
 - Allowing many programs to seemingly run at once
→ **Virtualizing the CPU**

Virtualizing Memory

- The physical memory is *an array of bytes*.
- A program keeps all of its data structures in memory.
 - **Read memory** (load):
 - Specify an address to be able to access the data
 - **Write memory** (store):
 - Specify the data to be written to the given address

Virtualizing Memory

- A program that Accesses Memory (mem.c)

```
1  #include <unistd.h>
2  #include <stdio.h>
3  #include <stdlib.h>
4  #include "common.h"
5
6  int
7  main(int argc, char *argv[])
8  {
9      int *p = malloc(sizeof(int)); // a1: allocate some memory
10     assert(p != NULL);
11     printf("(%) address of p: %08x\n",
12            getpid(), (unsigned) p); // a2: print out the
                                     address of the memory
13     *p = 0; // a3: put zero into the first slot of the memory
14     while (1) {
15         Spin(1);
16         *p = *p + 1;
17         printf("(%) p: %d\n", getpid(), *p); // a4
18     }
19     return 0;
20 }
```

Virtualizing Memory

- The output of the program `mem.c`

```
prompt> ./mem
(2134) memory address of p: 00200000
(2134) p: 1
(2134) p: 2
(2134) p: 3
(2134) p: 4
(2134) p: 5
^C
```

- The newly allocated memory is at address `00200000`.
- It updates the value and prints out the result.

Virtualizing Memory

- Running
mem.c
multiple times

```
prompt> ./mem &; ./mem &
[1] 24113
[2] 24114
(24113) memory address of p: 00200000
(24114) memory address of p: 00200000
(24113) p: 1
(24114) p: 1
(24114) p: 2
(24113) p: 2
(24113) p: 3
(24114) p: 3
...
```

- As if each running program has its **own private memory**.
 - Each running program has allocated memory at the same address.
 - Each seems to be updating the value at 00200000 independently.

Virtualizing Memory

- Each process accesses its own private **virtual address space**.
 - The OS maps **address space** onto the **physical memory**.
 - A memory reference within one running program does not affect the address space of other processes.
 - Physical memory is a **shared resource**, managed by the OS.

CONCURRENCY

The Problem of Concurrency

- Where does it stem from?
- The OS is juggling **many things at once**, first running one process, then another, and so forth.
- Modern **multi-threaded programs** also exhibit the concurrency problem.

The Problem of Concurrency

- A multi-threaded program (thread.c)

```
1    #include <stdio.h>
2    #include <stdlib.h>
3    #include "common.h"
4
5    volatile int counter = 0;
6    int loops;
7
8    void *worker(void *arg) {
9        int i;
10       for (i = 0; i < loops; i++) {
11           counter++;
12       }
13       return NULL;
14   }
15 ...
```

The Problem of Concurrency

- A multi-threaded program (thread.c)

```
16     int
17     main(int argc, char *argv[])
18     {
19         if (argc != 2) {
20             fprintf(stderr, "usage: threads <value>\n");
21             exit(1);
22         }
23         loops = atoi(argv[1]);
24         pthread_t p1, p2;
25         printf("Initial value : %d\n", counter);
26
27         Pthread_create(&p1, NULL, worker, NULL);
28         Pthread_create(&p2, NULL, worker, NULL);
29         Pthread_join(p1, NULL);
30         Pthread_join(p2, NULL);
31         printf("Final value : %d\n", counter);
32         return 0;
33     }
```

The Problem of Concurrency

- A multi-threaded program (thread.c)

```
16  int
17  main(int argc, char *argv[])
18  {
19      // The main program creates two threads.
20      // • Thread: runs within the same memory space. Each
21      //   thread starts running in a function called worker().
22      // • worker(): increments a counter
23      //
24      pthread_t p1, p2;
25      printf("Initial value : %d\n", counter);
26
27      Pthread_create(&p1, NULL, worker, NULL);
28      Pthread_create(&p2, NULL, worker, NULL);
29      Pthread_join(p1, NULL);
30      Pthread_join(p2, NULL);
31      printf("Final value : %d\n", counter);
32      return 0;
33  }
```

The Problem of Concurrency

- **loops** determines how many times each of the two workers will **increment the shared counter** in a loop.
 - loops: 1000.

```
prompt> gcc -o thread thread.c -Wall -pthread
prompt> ./thread 1000
Initial value : 0
Final value : 2000
```

- loops: 100000.

```
prompt> ./thread 100000
Initial value : 0
Final value : 143012 // huh??
prompt> ./thread 100000
Initial value : 0
Final value : 137298 // what the??
```

The Problem of Concurrency

- **Why is this happening?**
- Increment a shared counter → take three instructions.
 1. Load the value of the counter from memory into register.
 2. Increment it
 3. Store it back into memory
- These three instructions do not execute **atomically**. → Problem of **concurrency** happen.

PERSISTENCE

Persistence

- Devices such as DRAM store values in a volatile.
- *Hardware* and *software* are needed to store data **persistently**.
 - **Hardware:** I/O device such as a hard drive, solid-state drives(SSDs)
 - **Software:**
 - File system manages the disk.
 - File system is responsible for storing any files the user creates.

Persistence

- Create a file (`/tmp/file`) that contains the string "hello world"

```
1      #include <stdio.h>
2      #include <unistd.h>
3      #include <assert.h>
4      #include <fcntl.h>
5      #include <sys/types.h>
6
7      int
8      main(int argc, char *argv[])
9      {
10         int fd = open("/tmp/file", O_WRONLY | O_CREAT
                        | O_TRUNC, S_IRWXU);
11         assert(fd > -1);
12         int rc = write(fd, "hello world\n", 13);
13         assert(rc == 13);
14         close(fd);
15         return 0;
16     }
```

open(), **write()**, and **close()** system calls are routed to the part of OS called the file system, which handles the requests

Persistence

- What OS does in order to write to disk?
 - Figure out **where** on disk this new data will reside
 - **Issue I/O** requests to the underlying storage device
- File system handles system crashes during write.
 - **Journaling** or **copy-on-write**
 - Carefully ordering writes to disk

DESIGN GOALS OF AN OS

Design Goals

- Build up **abstraction**
 - Make the system convenient and easy to use.
- Provide high **performance**
 - Minimize the overhead of the OS.
 - OS must strive to provide virtualization without excessive overhead.
- **Protection** between applications
 - Isolation: Bad behavior of one does not harm other and the OS itself.

Design Goals

- High degree of **reliability**
 - The OS must also run non-stop.
- Other issues
 - Energy-efficiency
 - Security
 - Mobility

Reading Material

- **Chapter 2** of OSTEP book – by Remzi and Andrea Arpaci-Dusseau (University of Wisconsin)
<http://pages.cs.wisc.edu/~remzi/OSTEP/intro.pdf>

Questions?