

CIS 657 – Principles of Operating Systems

Topic: Process – Mechanism
(Limited Direct Execution)

Endadul Hoque

Acknowledgement


- Youjip Won (Hanyang University)
- OSTEP book – by Remzi and Andrea Arpaci-Dusseau (University of Wisconsin)

**Background: memory layout +
function calling**

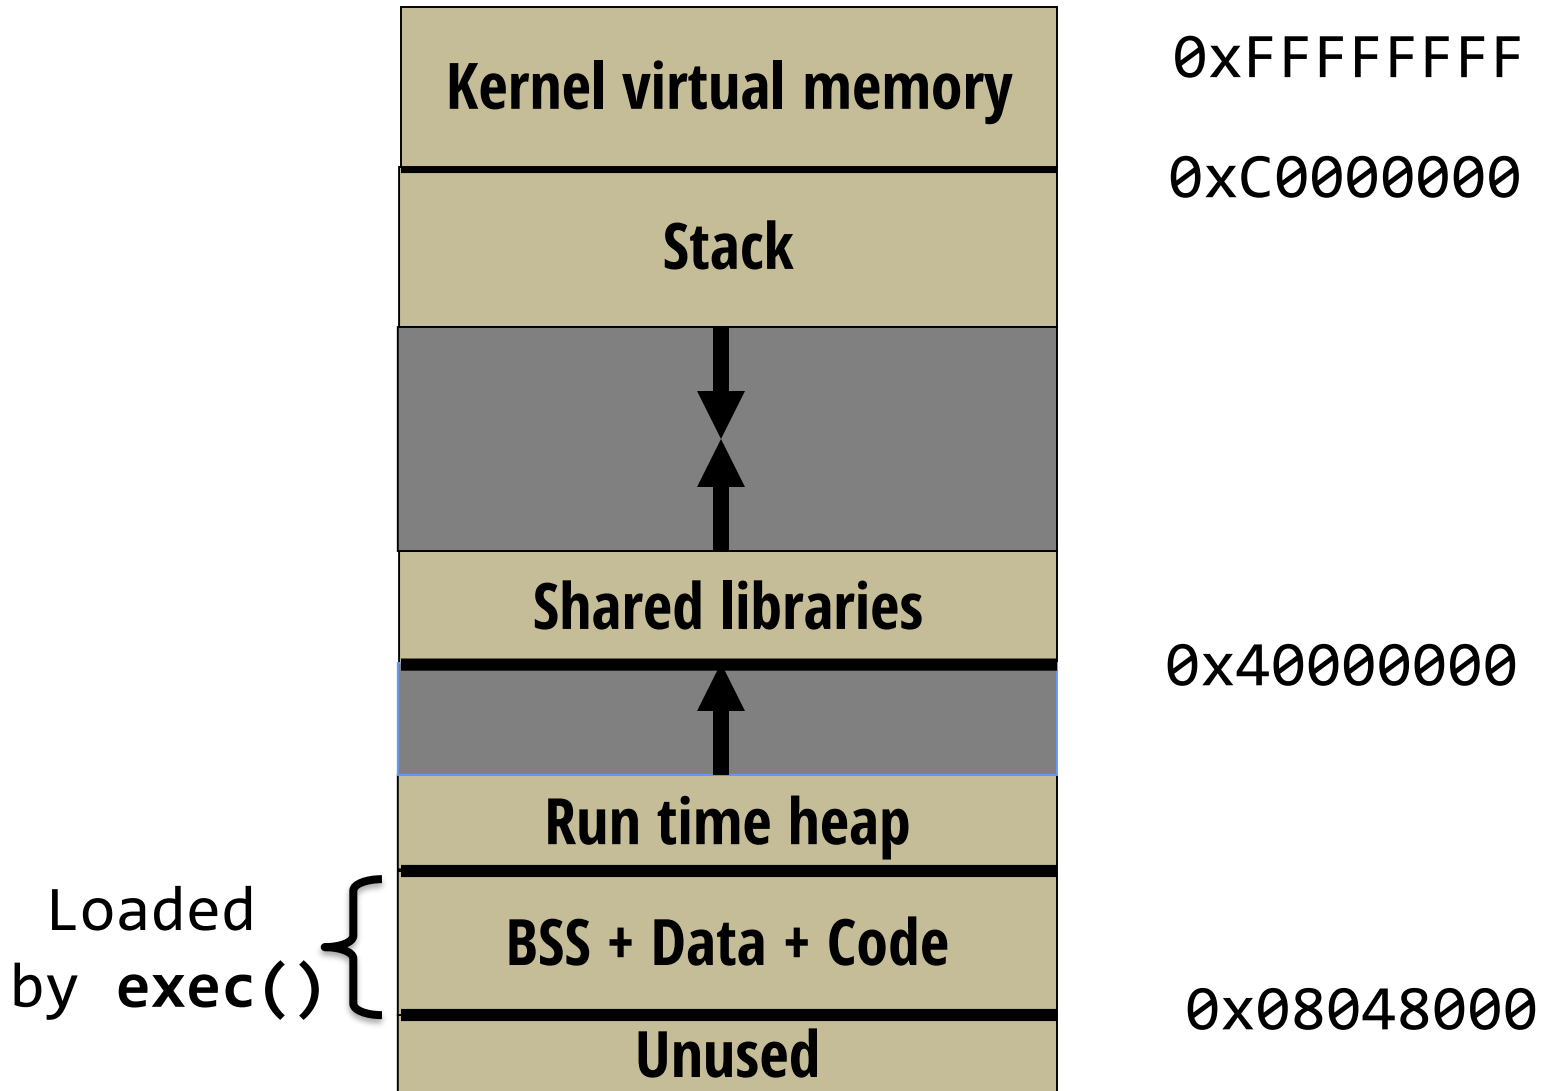
Revisit: Memory Layout of a Process

- The operating system creates a **process** by assigning memory and other resources
- **Code:** The program instructions to be executed
- **Data:** Initialized global and static variables
- **BSS:** un-initialized global and static variables (filled with 0 by OS)
- **Heap:** Dynamic memory for variables that are created with **malloc**, **calloc**, **realloc** and disposed of with **free**
- **Stack:** Keeps track of the point to which each active subroutine should **return** control when it finishes executing; **stores variables** that are local to functions

Virtual Memory



Revisit: Linux Process Memory Layout




Revisit: Execution of a C Program

- **PC (program counter** or instruction pointer) points to next machine instruction to be executed
- **Procedure call**
 - Prepare parameters
 - Save state (**SP** (stack frame pointer) and **PC**) and allocate on stack local variables
 - Here **PC** is pointing to the **return address**
 - Jumps to the beginning of procedure being called
- **Procedure return**
 - Recover state (**SP** and **PC** (this is return address)) from stack and adjust stack
 - Execution continues from **return address**


Revisit: Stack Frame

- **Parameters** for the procedure
- Save current PC onto stack (**return address**)
- Save current **SP (Stack Frame Pointer)** value onto stack
- Allocates stack space for **local variables** by decrementing SP by appropriate amount



Revisit: Stack Frame

- **Caller** : push parameter(s), push ret addr, push stack frame pointer, allocate space for locals, jmp to start addr of the callee
- **Callee**: execute the instructions, restore stack frame pointer
- **Return**: pop ret addr, jmp to addr



LIMITED DIRECT EXECUTION

How to efficiently virtualize the CPU?

- The OS needs to share the physical CPU by **time sharing**
 - Running one process, then stopping it and running another
- Issues
 - **Performance**: How can we implement virtualization without adding excessive overhead to the system?
 - **Control**: How can we run processes efficiently while retaining control over the CPU?

Direct Execution

- Just run the program directly on the CPU

OS	Program
<ol style="list-style-type: none">1. Create entry for process list2. Allocate memory for program3. Load program into memory4. Set up stack with argc / argv5. Clear registers6. Execute call main() <ol style="list-style-type: none">9. Free memory of process10. Remove from process list	<ol style="list-style-type: none">7. Run main()8. Execute return from main()

Direct Execution

- There are a few problems to *direct execution* approach
 - How can OS ensure that the program does not do anything unexpected?
 - How can OS take control of the CPU from the running process?

Without **limits** on running programs,
the OS wouldn't be in control of anything and
thus would be “**just a library**”

Problem 1: Restricted Operation

- What if a process wishes to perform some kind of restricted operation such as
 - Issuing an I/O request to a disk
 - Gaining access to more system resources such as CPU or memory
- **Solution:** Using **protected control transfer**
 - **User mode:** Applications do not have full access to hardware resources.
 - **Kernel mode:** The OS has access to the full resources of the machine

System Calls

- Allow the kernel to **carefully expose** certain key pieces of functionality to user program, e.g.,
 - Accessing the file system
 - Creating and destroying processes
 - Communicating with other processes
 - Allocating more memory

System Calls

- To execute a system call, program executes
 - **Trap** instruction
 - Jump into the kernel
 - Raise the privilege level to kernel mode
- When finished, the OS executes
 - **Return-from-trap** instruction
 - Return into the calling user program
 - Reduce the privilege level back to user mode

System Calls

- Hardware involvement
 - Trap
 - The processor pushes the process's register context onto a per-process **kernel stack**
 - Return-from-trap
 - Pop the values off the stack and resume execution

System Calls

How does the trap know which code to run inside the OS?

- Does the calling process specify where to jump?
 - No, but **why not?**
- The kernel must carefully control what code executes upon a trap
- Solution: **Trap table**
 - Specify the **location** of the trap handlers
 - what the code to jump to and execute when a *system call* or other *exceptional events* occur (e.g., interrupt)
 - At boot time, OS initializes the trap table and informs the hardware (also a *privileged* operation)

System Calls

- To specify the **exact system call**
 - User code cannot specify the exact address to jump to
 - Instead, it uses a **system-call number** ; that is a unique number for each system call
- Execution
 - User code places the desired syscall number to a specific *register* or memory *location*
 - OS reads the number inside trap handler
 - If valid, executes the corresponding code
- This level of indirection provides a **protection**

Limited Direct Execution Protocol (1/2)

OS @ boot
(kernel mode)

Hardware

initialize trap table

remember address of ...
syscall handler

OS @ run
(kernel mode)

Hardware

Program
(user mode)

Create entry for process list
Allocate memory for program
Load program into memory
Setup user stack with argv
Fill kernel stack with reg/PC
return-from -trap

restore regs from kernel stack
move to user mode
jump to main

Run main()
...
Call system call
trap into OS

Limited Direct Execution Protocol (2/2)

OS @ run
(kernel mode)

Hardware

Program
(user mode)

(Cont.)

Handle trap
Do work of syscall
return-from-trap

save regs to kernel stack
move to kernel mode
jump to trap handler

restore regs from kernel stack
move to user mode
jump to PC after trap

Free memory of process
Remove from process list

...
return from main
trap (via **exit()**)

Problem 2: Switching Between Processes

- How can the OS **regain control** of the CPU so that it can switch between *processes*?
 - A cooperative Approach: **Wait for system calls**
 - A Non-Cooperative Approach: **The OS takes control**

A cooperative Approach: Wait for system calls

- Processes **periodically give up the CPU** by making **system calls** such as `yield`.
 - The OS decides to run some other task.
 - Application also transfer control to the OS when they do something illegal.
 - Divide by zero
 - Try to access memory that it shouldn't be able to access
 - E.g., Early versions of the Macintosh OS, The old Xerox Alto system

A process gets stuck in an infinite loop.
→ **Reboot the machine**

A Non-Cooperative Approach: OS Takes Control

- **A timer interrupt**

- During the boot sequence, the OS start the timer.
- The timer raise an interrupt every so many milliseconds.
- When the interrupt is raised :
 - The currently running process is halted.
 - Save enough of the state of the program
 - A pre-configured interrupt handler in the OS runs.

A **timer interrupt** gives OS the ability to run again on a CPU.

Saving and Restoring Context

- **Scheduler** makes a decision:
 - Whether to continue running the **current process**, or switch to a **different one**.
 - If the decision is made to switch, the OS executes context switch.

Context Switch

- A low-level piece of assembly code
 - **Save a few register values** for the current process onto its kernel stack
 - General purpose registers
 - PC
 - kernel stack pointer
 - **Restore registers** for the soon-to-be-executing process from its kernel stack
 - **Switch to the kernel stack** for the soon-to-be-executing process

Limited Direction Execution Protocol (with Timer interrupt)

OS @ boot
(kernel mode)

Hardware

initialize trap table

remember address of ...
syscall handler
timer handler

start interrupt timer

start timer
interrupt CPU in X ms

OS @ run
(kernel mode)

Hardware

Program
(user mode)

Process A

...

timer interrupt

save regs(A) to k-stack(A)
move to kernel mode
jump to trap handler

Limited Direction Execution Protocol (with Timer interrupt)

OS @ run
(kernel mode)

Hardware

Program
(user mode)

(Cont.)

Handle the trap

Call switch() routine

save regs(A) to proc-struct(A)

restore regs(B) from proc-struct(B)

switch to k-stack(B)

return-from-trap (into B)

restore regs(B) from k-stack(B)

move to user mode

jump to B's PC

Process B

...

The xv6 Context Switch Code

```
1 # void swtch(struct context **old, struct context *new);
2 #
3 # Save current register context in old
4 # and then load register context from new.
5 .globl swtch
6 swtch:
7     # Save old registers
8     movl 4(%esp), %eax           # put old ptr into eax
9     popl 0(%eax)                # save the old IP
10    movl %esp, 4(%eax)           # and stack
11    movl %ebx, 8(%eax)           # and other registers
12    movl %ecx, 12(%eax)
13    movl %edx, 16(%eax)
14    movl %esi, 20(%eax)
15    movl %edi, 24(%eax)
16    movl %ebp, 28(%eax)
17
```

The xv6 Context Switch Code

```
1 # void swtch(struct context **old, struct context *new);
2 #
3 # Save current register context in old
4 # and then load register context from new.
```

```
18      # Load new registers
19      movl 4(%esp), %eax # put new ptr into eax
20      movl 28(%eax), %ebp # restore other registers
21      movl 24(%eax), %edi
22      movl 20(%eax), %esi
23      movl 16(%eax), %edx
24      movl 12(%eax), %ecx
25      movl 8(%eax), %ebx
26      movl 4(%eax), %esp # stack is switched here
27      pushl 0(%eax)      # return addr put in place
28      ret                # finally return into new ctxt
```

Worried About Concurrency?

- What happens if, during interrupt or trap handling, another interrupt occurs?
- OS handles these situations:
 - **Disable interrupts** during interrupt processing
 - Use a number of sophisticated **locking** schemes to protect concurrent access to internal data structures.

Reading Material

- **Chapter 6** of OSTEP book – by Remzi and Andrea Arpaci-Dusseau (University of Wisconsin)
<http://pages.cs.wisc.edu/~remzi/OSTEP/cpu-mechanisms.pdf>

Questions?