

CIS 657 – Principles of Operating Systems

Topic: Process – Policies
(Scheduling)

Endadul Hoque

Acknowledgement

- Youjip Won (Hanyang University)
- OSTEP book – by Remzi and Andrea Arpaci-Dusseau (University of Wisconsin)

Process Scheduling

- Workload assumptions:
 1. Each job runs for the **same amount of time**.
 2. All jobs **arrive** at the same time.
 3. All jobs only use the **CPU** (i.e., they perform no I/O).
 4. The **run-time** of each job is known.

Process Scheduling


- Performance metric: **Turnaround time**
 - The time at which **the job completes** minus the time at which **the job arrived** in the system.

$$T_{turnaround} = T_{completion} - T_{arrival}$$

- Another metric is **fairness**.
 - Performance and fairness are often **at odds** in scheduling.


First in, First Out (FIFO)

- Also, known as First Come, First Served (**FCFS**)
 - Very simple and easy to implement
- Example:
 - A arrived just before B which arrived just before C.
 - Each job runs for 10 seconds.




$$\text{Average turnaround time} = \frac{10 + 20 + 30}{3} = 20 \text{ sec}$$

Why is FIFO not that great? – Convoy effect



Why is FIFO not that great? – Convoy effect


- Let's relax assumption 1: Each job **no longer** runs for the same amount of time.
- Example:
 - A arrived just before B which arrived just before C.
 - A runs for 100 seconds, B and C run for 10 each.



$$\text{Average turnaround time} = \frac{100 + 110 + 120}{3} = 110 \text{ sec}$$

Shortest Job First (SJF)


- Run the shortest job first, then the next shortest, and so on
 - Non-preemptive scheduler
- Example:
 - A arrived just before B which arrived just before C.
 - A runs for 100 seconds, B and C run for 10 each.



$$\text{Average turnaround time} = \frac{10 + 20 + 120}{3} = 50 \text{ sec}$$

SJF with Late Arrivals from B and C

- Let's relax assumption 2: Jobs can arrive at any time.
- Example:
 - A arrives at $t=0$ and needs to run for 100 seconds.
 - B and C arrive at $t=10$ and each need to run for 10 seconds




$$\begin{aligned} \text{Average turnaround time} &= \frac{100 + (110 - 10) + (120 - 10)}{3} \\ &= 103.33 \text{ sec} \end{aligned}$$

Shortest Time-to-Completion First (STCF)

- Add **preemption** to SJF
 - Also known as Preemptive Shortest Job First (PSJF)
- A new job enters the system:
 - Determine of the remaining jobs and new job
 - Schedule the job which has the lowest time left (i.e., which requires the minimum time to finish)

Shortest Time-to-Completion First (STCF)

- Example:
 - A arrives at $t=0$ and needs to run for 100 seconds.
 - B and C arrive at $t=10$ and each need to run for 10 seconds



$$\begin{aligned} \text{Average turnaround time} &= \frac{(120 - 0) + (20 - 10) + (30 - 10)}{3} \\ &= 50 \text{ sec} \end{aligned}$$

New Scheduling Metric: Response Time

- The time from **when the job arrives** to the **first time it is scheduled**.

$$T_{response} = T_{firstrun} - T_{arrival}$$

- STCF and related disciplines are not particularly good for response time.

How can we build a scheduler that is
sensitive to response time?


Round Robin (RR) Scheduling

- Time slicing Scheduling
 - Run a job for a **time slice** and then switch to the next job in the **run queue** until the jobs are finished.
 - Time slice is sometimes called a scheduling quantum.
 - It repeatedly does so until the jobs are finished.
 - The length of a time slice must be *a multiple of* the timer-interrupt period.

RR is fair, but performs poorly on metrics such as turnaround time

Round Robin (RR) Scheduling


- A, B and C arrive at the same time.
- They each wish to run for 5 seconds.



SJF (Bad for Response Time)

$$T_{average\ response} = \frac{0 + 5 + 10}{3} = \text{5sec}$$

ABCABCABCABCABC



RR with a time-slice of 1sec (Good for Response Time)

$$T_{average\ response} = \frac{0 + 1 + 2}{3} = \text{1sec}$$

The length of the time slice is critical


- The shorter time slice
 - Better response time
 - The cost of context switching will dominate overall performance.
- The longer time slice
 - Amortize the cost of switching
 - Worse response time

Deciding on the length of the time slice presents a **trade-off** to a system designer


Incorporating I/O

- Let's relax assumption 3: All programs perform I/O
- Example:
 - A and B need 50ms of CPU time each.
 - A runs for 10ms and then issues an I/O request
 - I/Os each take 10ms
 - B simply uses the CPU for 50ms and performs no I/O
 - The scheduler runs A first, then B after

Incorporating I/O



Poor Use of Resources



Overlap Allows Better Use of Resources

Maximize the
CPU utilization

Incorporating I/O

- When a job initiates an I/O request.
 - The job is blocked waiting for I/O completion.
 - The scheduler should schedule another job on the CPU.
- When the I/O completes
 - An interrupt is raised.
 - The OS moves the process from blocked back to the ready state.

Reading Material

- **Chapter 7** of OSTEP book – by Remzi and Andrea Arpaci-Dusseau (University of Wisconsin)
<http://pages.cs.wisc.edu/~remzi/OSTEP/cpu-sched.pdf>

Questions?