

CIS 657 – Principles of Operating Systems

Topic: Memory – Address Translation

Endadul Hoque

Acknowledgement

- Youjip Won (Hanyang University)
- OSTEP book – by Remzi and Andrea Arpaci-Dusseau (University of Wisconsin)

Memory Virtualizing with Efficiency and Control

- How to build an efficient virtualization of memory?
- How to maintain control over memory to enable protection?
- How to provide the flexibility needed by applications?

Memory Virtualizing with Efficiency and Control

- Memory virtualizing takes a similar strategy known as **limited direct execution(LDE)** for efficiency and control.
- In memory virtualizing, efficiency and control are attained by **hardware support**.
 - e.g., registers, TLB(Translation Look-aside Buffer)s, page-table

Address Translation

- Hardware transforms a **virtual address** to a **physical address**.
 - The desired information is actually stored in a physical address.
- The OS must get involved at key points to set up the hardware.
 - The OS must manage memory to judiciously intervene.

Example: Address Translation

- C - Language code

```
void func(){  
    int x;  
    ...  
    x = x + 3; // this is the line of code we are interested in
```

- **Load** a value from memory
- **Increment** it by three
- **Store** the value back into memory

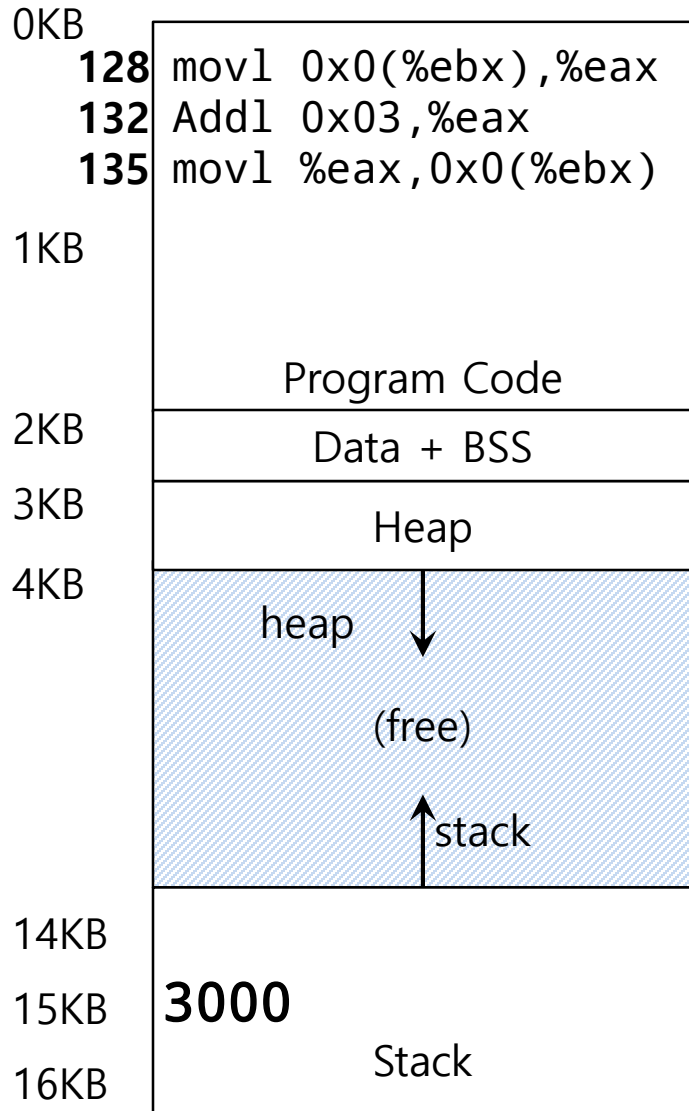
Example: Address Translation

- Assembly

```
128 : movl 0x0(%ebx), %eax ; load 0+ebx into eax
132 : addl $0x03, %eax      ; add 3 to eax register
135 : movl %eax, 0x0(%ebx)  ; store eax back to mem
```

- Presume that the address of 'x' has been place in `ebx` register.
- **Load** the value at that address into `eax` register.
- **Add** 3 to `eax` register.
- **Store** the value in `eax` back into memory.

Example: Address Translation



- Fetch instruction at address **128**
- Execute this instruction (load from address 15KB)
- Fetch instruction at address 132
- Execute this instruction (no memory reference)
- Fetch the instruction at address 135
- Execute this instruction (store to address 15 KB)

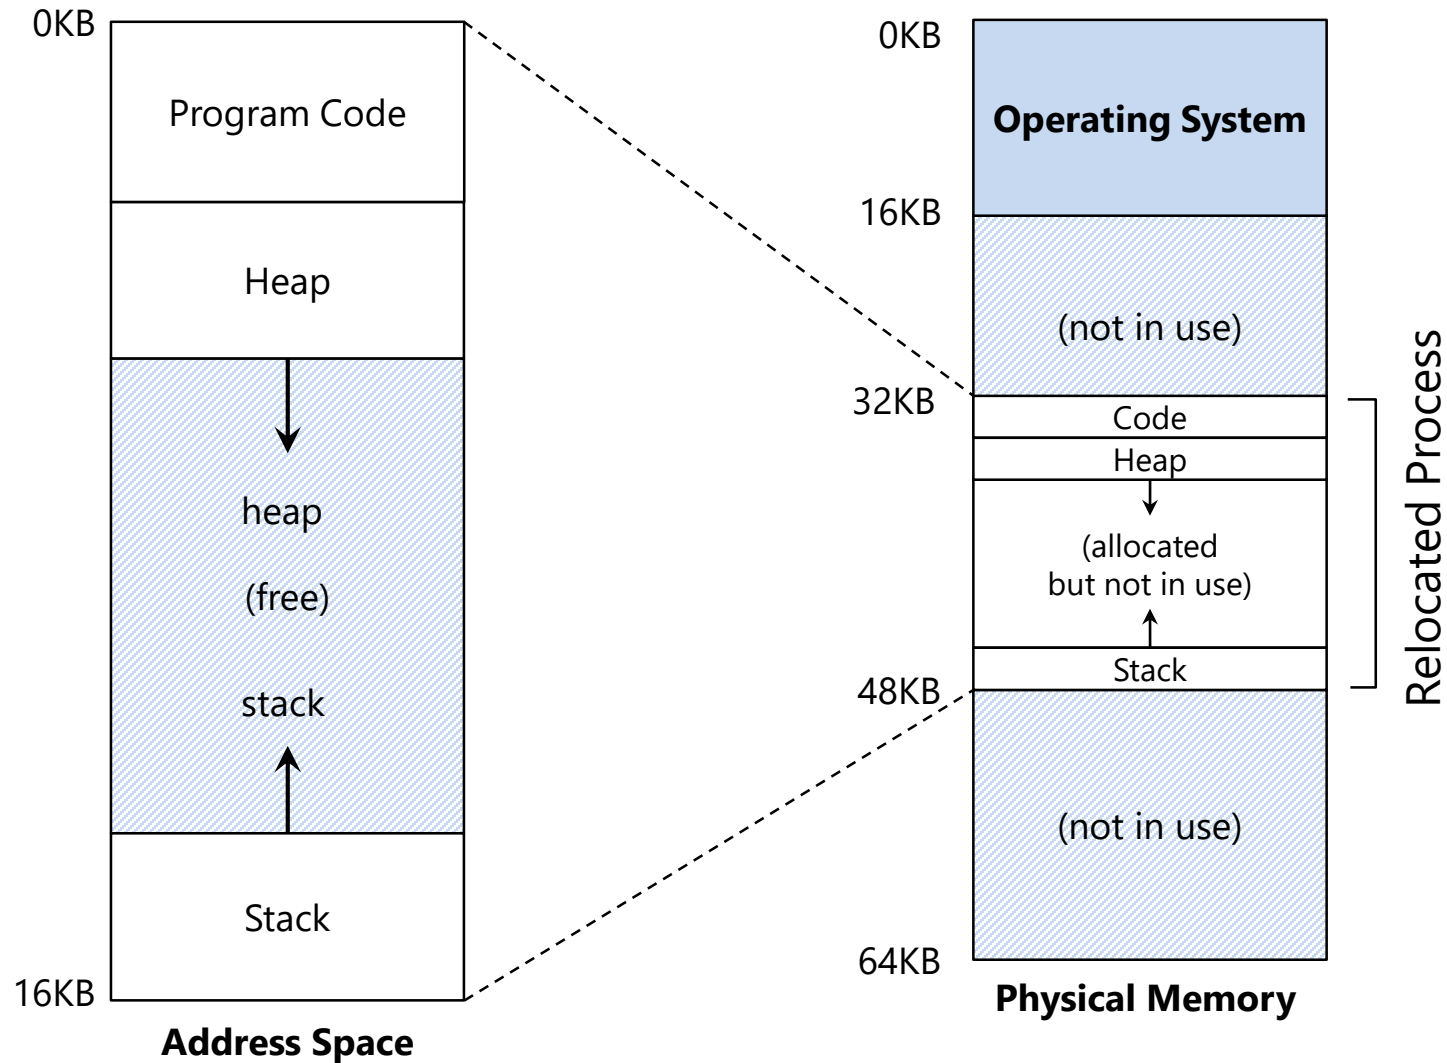
Relocation Address Space

- The OS wants to place the process **somewhere else** in physical memory, not at address 0.
 - The address space start at address 0.

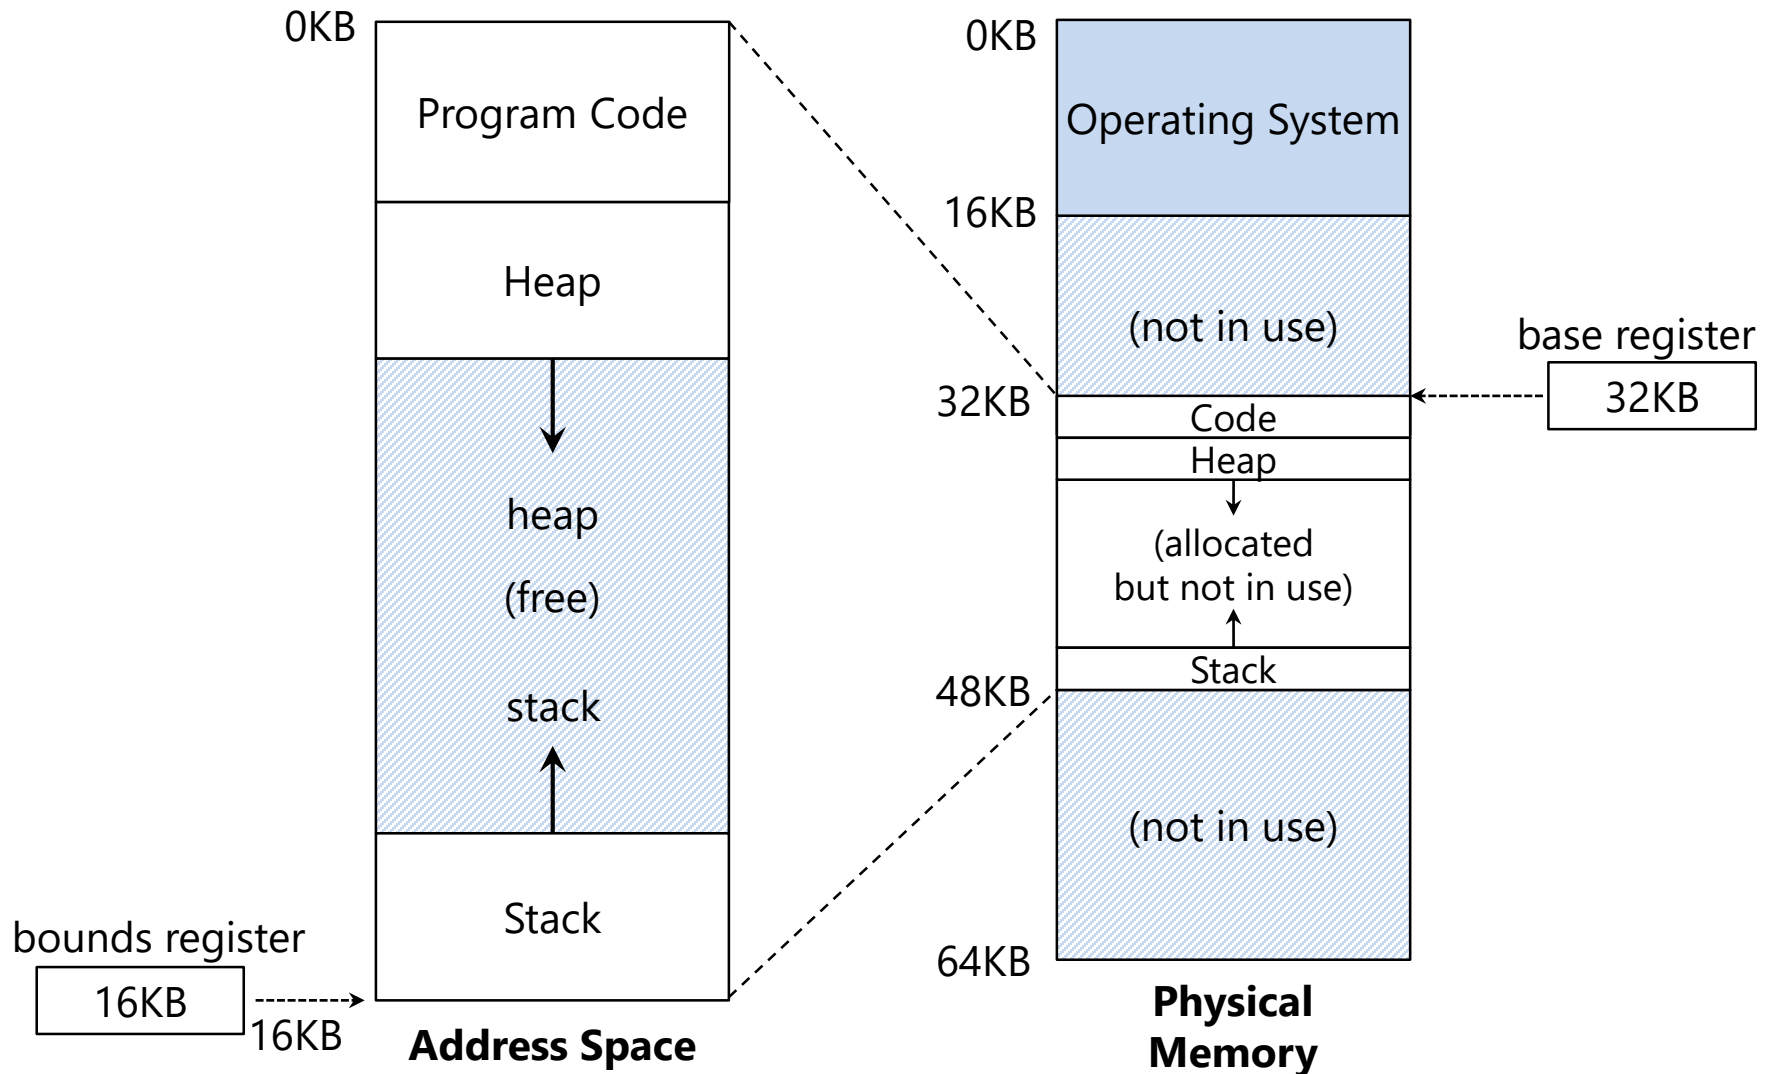
Working Assumptions

1. The address space must be placed **contiguously** in physical memory
2. The size of the address space is **less than** the size of physical memory
3. Each address space is exactly the **same size**

A Single Relocated Process



Base and Bounds Register



Dynamic (Hardware base) Relocation

- When a program starts running, the OS decides **where** in physical memory a process should be **loaded**.
 - Set the **base** register a value.

$$\text{physical address} = \text{virtual address} + \text{base}$$

- Every virtual address must **not be greater than bound** and **negative**.

$$0 \leq \text{virtual address} < \text{bounds}$$

Relocation and Address Translation


128 : movl 0x0(%ebx), %eax

- **Fetch** instruction at address 128

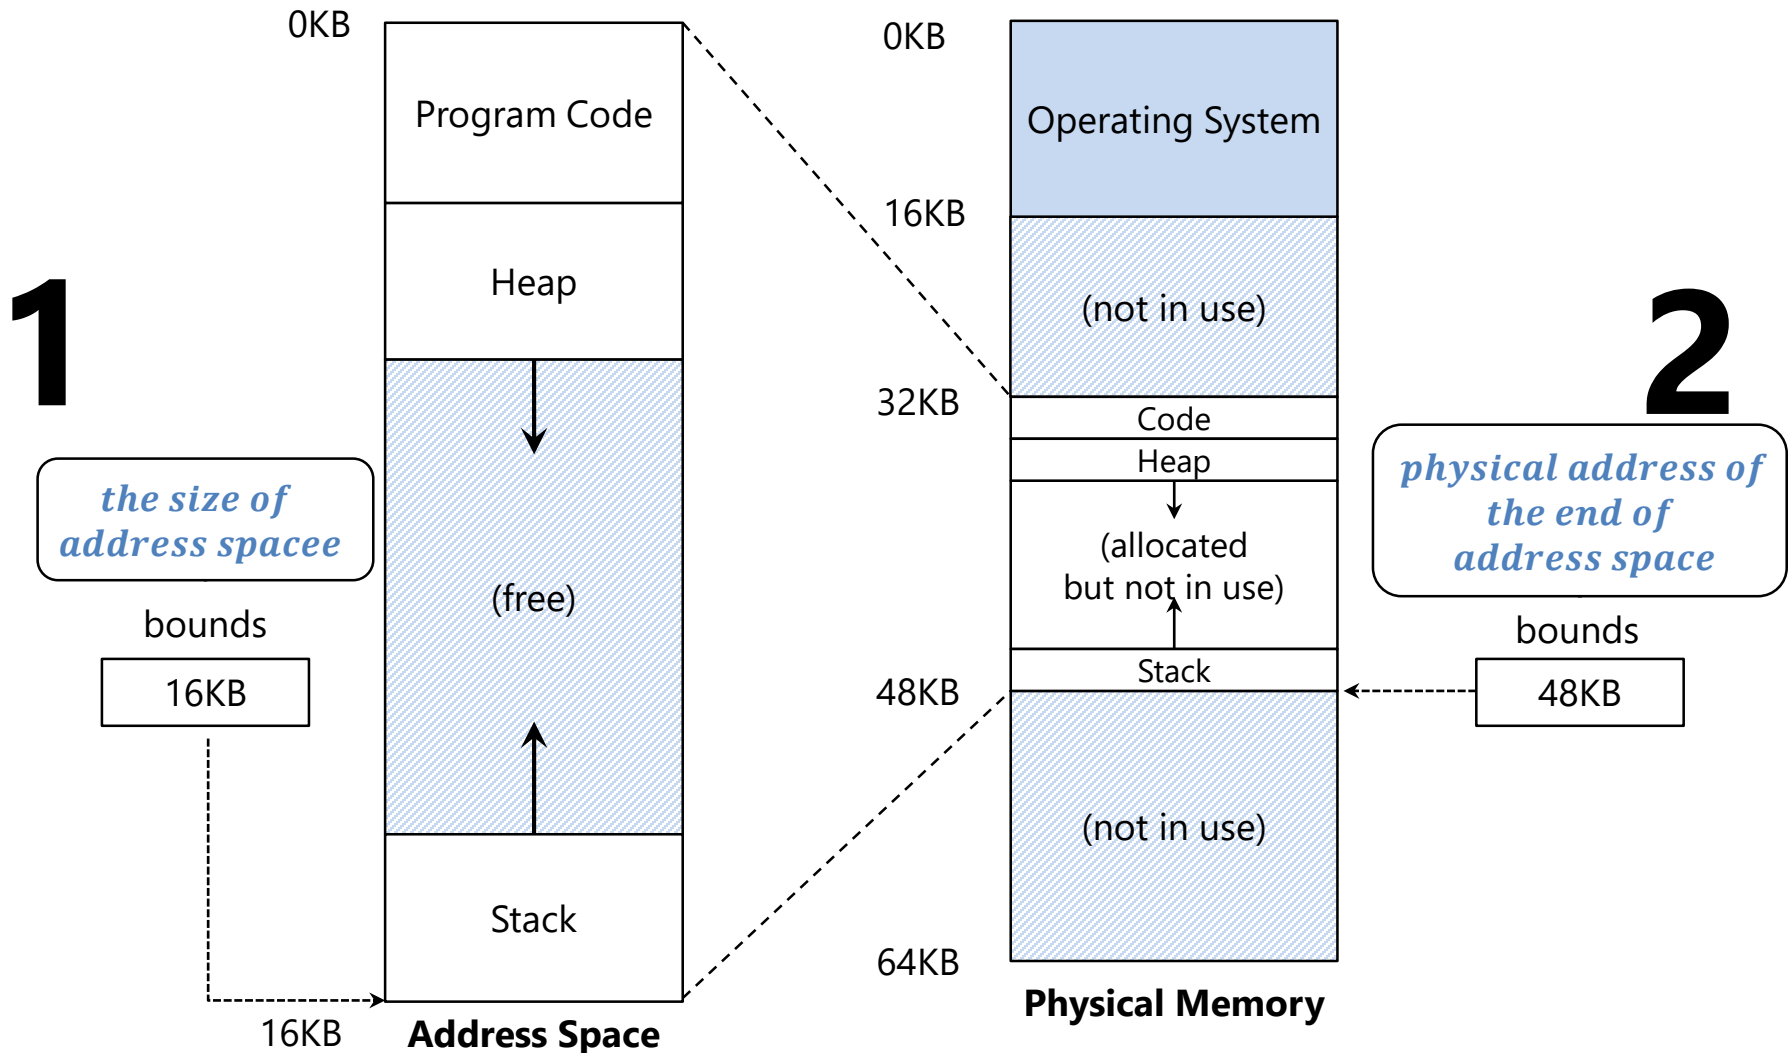
$$32896 = 128 + 32KB(base)$$

- **Execute** this instruction
 - Load from address 15KB

$$47KB = 15KB + 32KB(base)$$



Two ways to use Bounds Register



OS Issues for Memory Virtualizing



- The OS must **take action** to implement **base-and-bounds** approach.
- Three critical junctures:
 - When a process **starts running**:
 - Finding space for address space in physical memory
 - When a process is **terminated**:
 - Reclaiming the memory for use
 - When context **switch occurs**:
 - Saving and storing the base-and-bounds pair

OS Issues: When a Process Starts Running

- The OS must **find a room** for a new address space.
 - **free list** : A list of the range of the physical memory which are not in use.



The OS lookup the free list

Free list



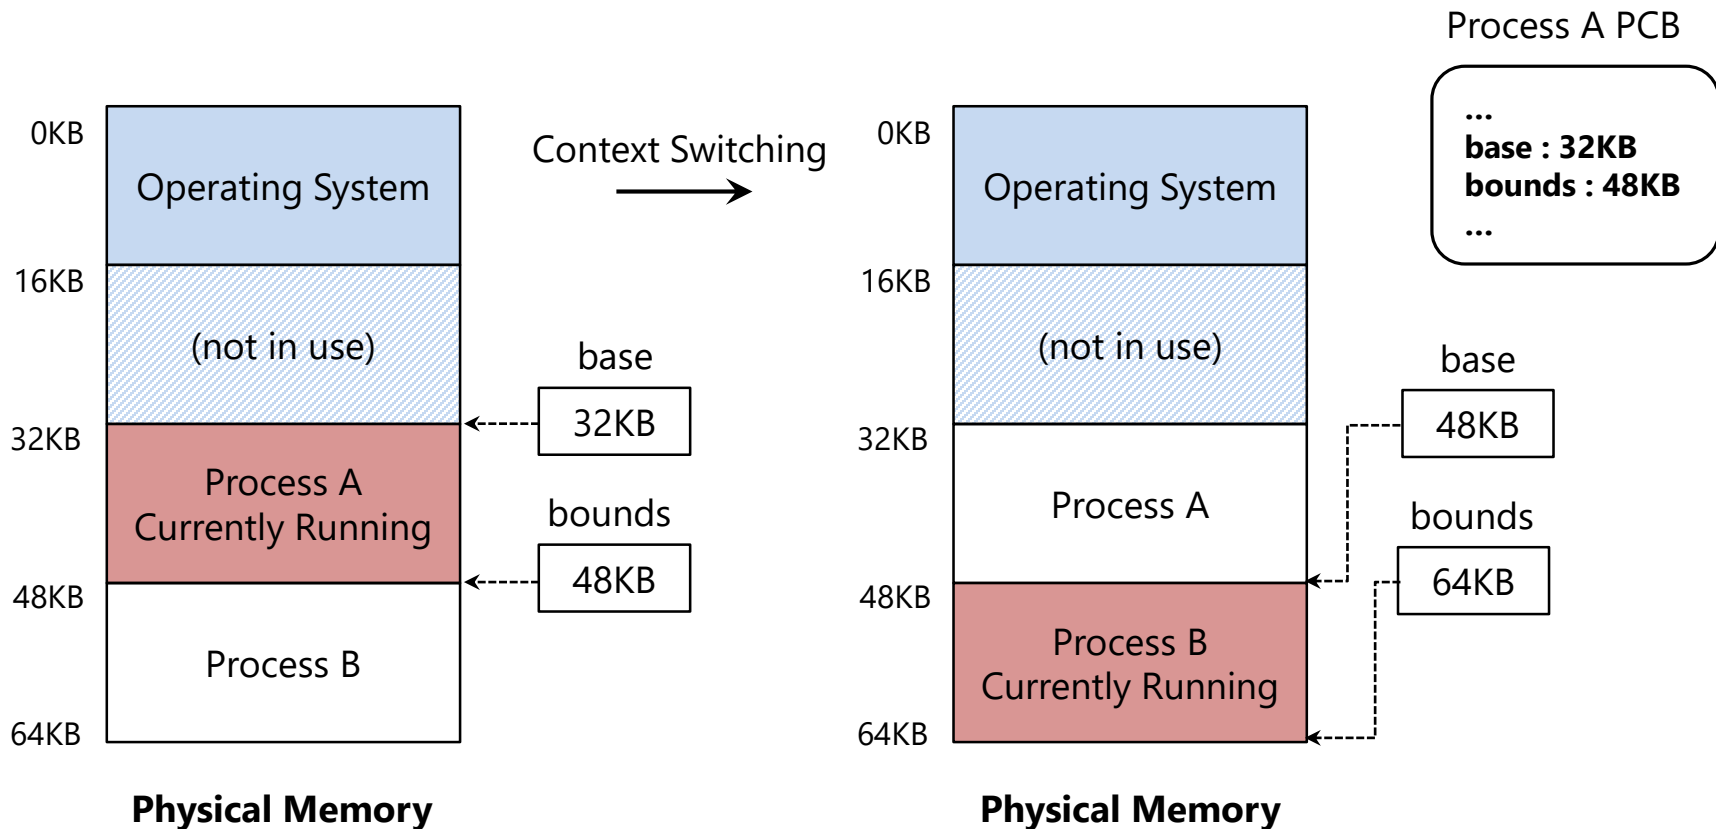
OS Issues: When a Process Is Terminated

- The OS must **put the memory back** on the free list.



OS Issues: When Context Switch Occurs

- The OS must **save and restore** the base-and-bounds pair.
 - In **process structure** or **process control block (PCB)**



Hardware support

- Hardware must provide the base and bounds registers
 - A **pair of registers per CPU** to support address translation and bounds check
- Part of the memory management unit (**MMU**) of the CPU
- Hardware should provide **special instructions** to modify these registers
 - OS needs to change them when different processes run
- CPU must be able to **generate/raise exceptions** in case of illegal memory access
 - CPU must stop executing the current process and run the OS “out-of-bound” exception handler

Limited Direct Execution (w/ Dynamic Relocation)

OS @ boot
(kernel mode)

Hardware

(No Program Yet)

initialize trap table

remember addresses of...
system call handler
timer handler
illegal mem-access handler
illegal instruction handler

start interrupt timer

start timer; interrupt after X ms

initialize process table
initialize free list

@Boot

OS @ run (kernel mode)	Hardware	Program (user mode)
To start process A: allocate entry in process table alloc memory for process set base/bound registers return-from-trap (into A)	restore registers of A move to user mode jump to A's (initial) PC	
	translate virtual address perform fetch	Process A runs Fetch instruction
	if explicit load/store: ensure address is legal translate virtual address perform load/store	Execute instruction
	Timer interrupt move to kernel mode jump to handler	(A runs...)

Handle timer

decide: stop A, run B
call `switch()` routine
 save `regs(A)`
 to `proc-struct(A)`
 (including base/bounds)
 restore `regs(B)`
 from `proc-struct(B)`
 (including base/bounds)
return-from-trap (into B)

restore registers of B
move to **user mode**
jump to B's PC

Process B runs
Execute bad load

Load is out-of-bounds;
move to **kernel mode**
jump to trap handler

Handle the trap

decide to kill process B
deallocate B's memory
free B's entry
 in process table

Reading Material

- **Chapter 15** of OSTEP book – by Remzi and Andrea Arpaci-Dusseau (University of Wisconsin)
<http://pages.cs.wisc.edu/~remzi/OSTEP/vm-mechanism.pdf>

Questions?