

CIS 657 – Principles of Operating Systems

Topic: Persistence – I/O Devices

Endadul Hoque


Acknowledgement

- Youjip Won (Hanyang University)
- OSTEP book – by Remzi and Andrea Arpaci-Dusseau (University of Wisconsin)

I/O Devices

- I/O is **critical** to computer system to **interact with systems**.
- Issue :
 - How should I/O be integrated into systems?
 - What are the general mechanisms?
 - How can we make them efficient?

System Architecture



Prototypical System Architecture

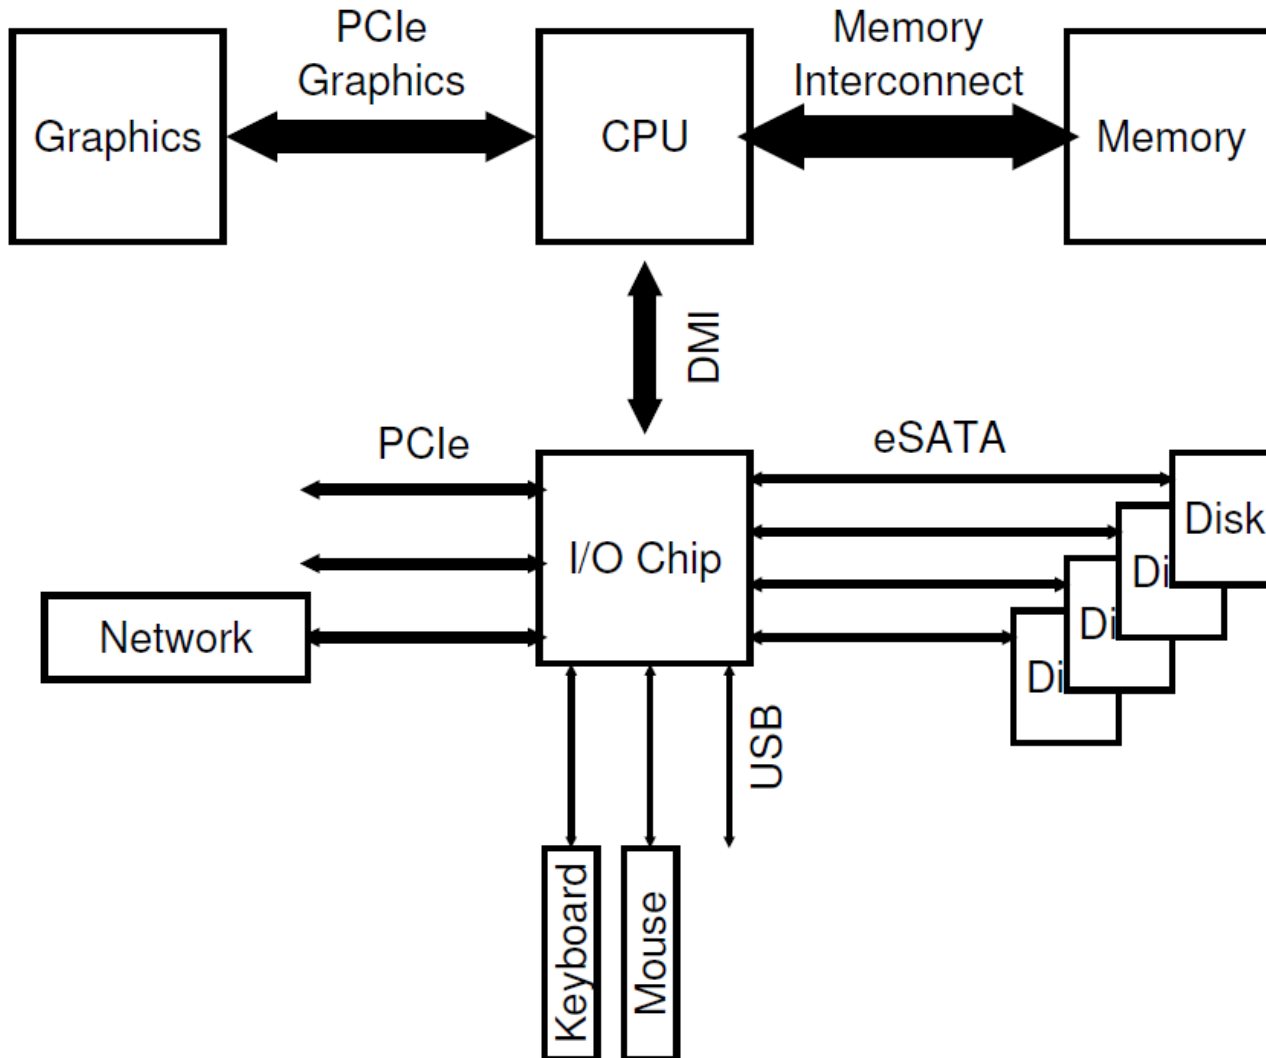
CPU is attached to the main memory of the system via some kind of memory **bus.**

Some devices are connected to the system via a general **I/O bus.**

System Architecture

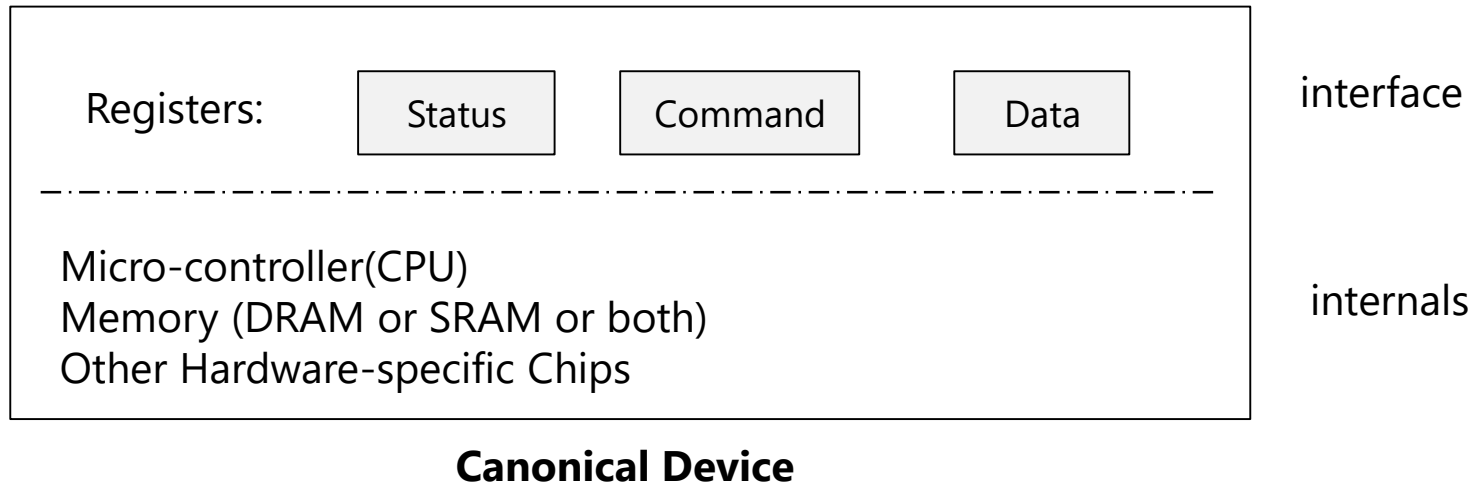
- Buses
 - Data paths that are provided to exchange information between CPU(s), RAM, and I/O devices.
- I/O bus
 - Data path that connects a CPU to an I/O device.
 - I/O bus is connected to I/O device by hardware components: I/O ports, interfaces and device controllers.

Modern System Architecture



Canonical Device

- Canonical Devices has two important components.
 - **Hardware interface** allows the system software to control its operation.
 - **Internals** which is implementation specific.



Hardware interface of Canonical Device

- **status register**
 - See the current status of the device
- **command register**
 - Tell the device to perform a certain task
- **data register**
 - Pass data to the device, or get data from the device

**By reading and writing above three registers,
the operating system can control device behavior.**

Hardware interface of Canonical Device

- Typical interaction example

```
while ( STATUS == BUSY)
    ; //wait until device is not busy
write data to data register
write command to command register
    Doing so starts the device and executes the command
while ( STATUS == BUSY)
    ; //wait until device is done with your request
```

Polling

- Operating system waits until the device is ready by **repeatedly** reading the status register.
 - Positive aspect is simple and working.
 - **However, it wastes CPU time just waiting for the device.**
 - Switching to another ready process is better utilizing the CPU.

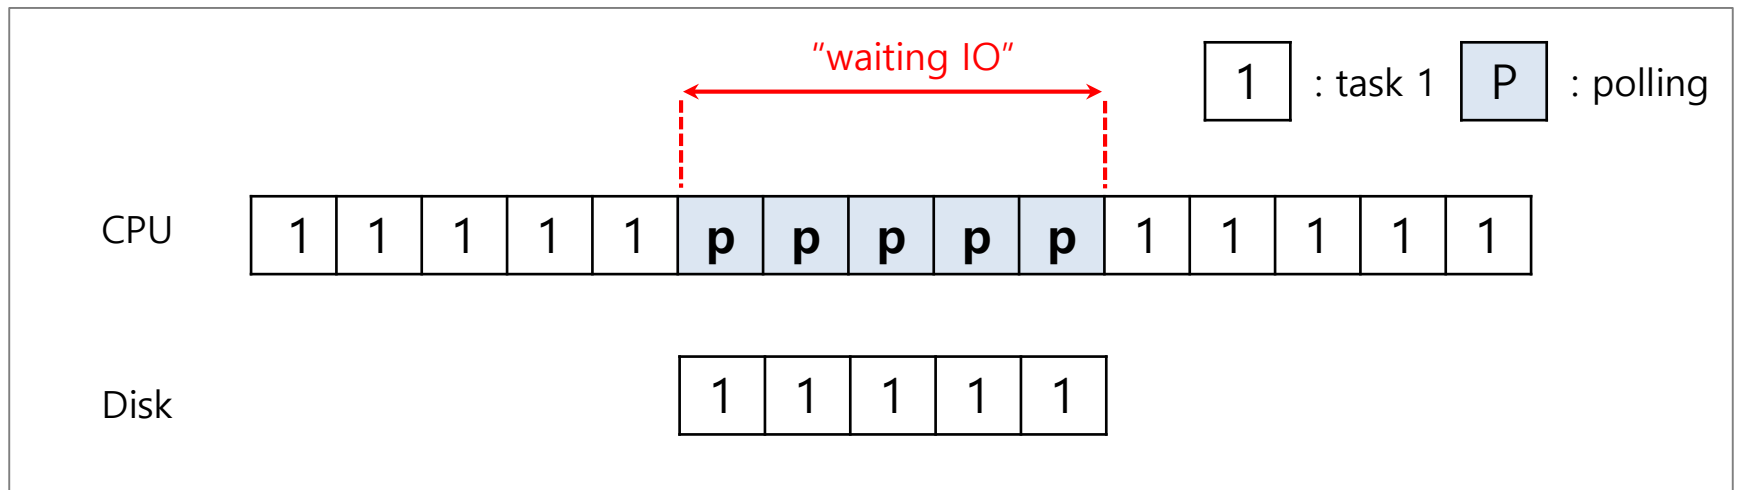


Diagram of CPU utilization by polling

Interrupts

- **Put the I/O request process to sleep** and context switch to another.
- When the device is finished, wake the process waiting for the I/O by **interrupt**.
 - Positive aspect : **CPU and the disk are properly utilized.**




Diagram of CPU utilization by interrupt

Polling vs. Interrupts

- *However, “interrupts is not always the best solution”*
 - If, device performs very quickly, interrupt will “slow down” the system.
 - Because **context switch is expensive (switching to another process)**

If a device is fast → **poll** is best.
If it is slow → **interrupt** is better.

CPU is once again over-burdened

- CPU **wastes a lot of time** to copy the *a large chunk of data* from memory to the device.

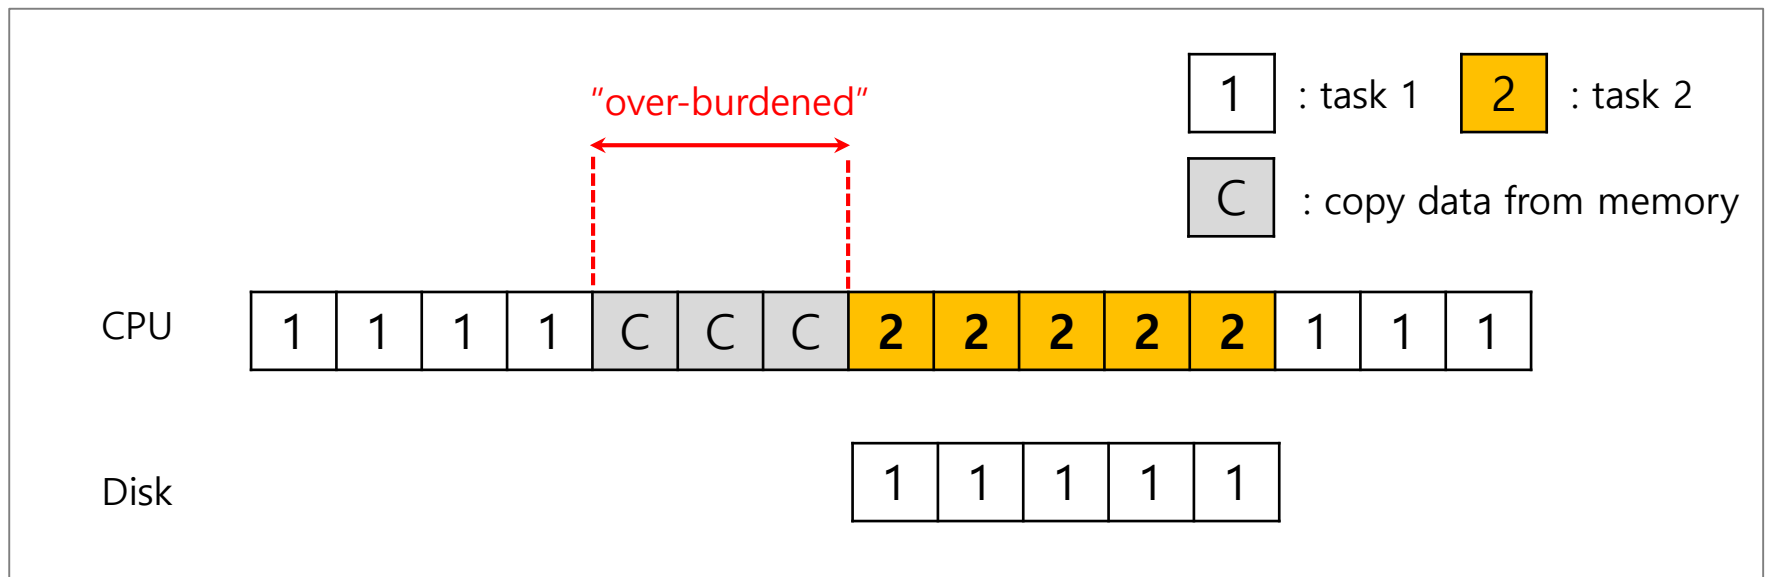


Diagram of CPU utilization

DMA (Direct Memory Access)

- **Copy data** in memory by knowing “where the data lives in memory, how much data to copy”
- When completed, DMA raises an interrupt, I/O begins on Disk.




Diagram of CPU utilization by DMA

Device interaction

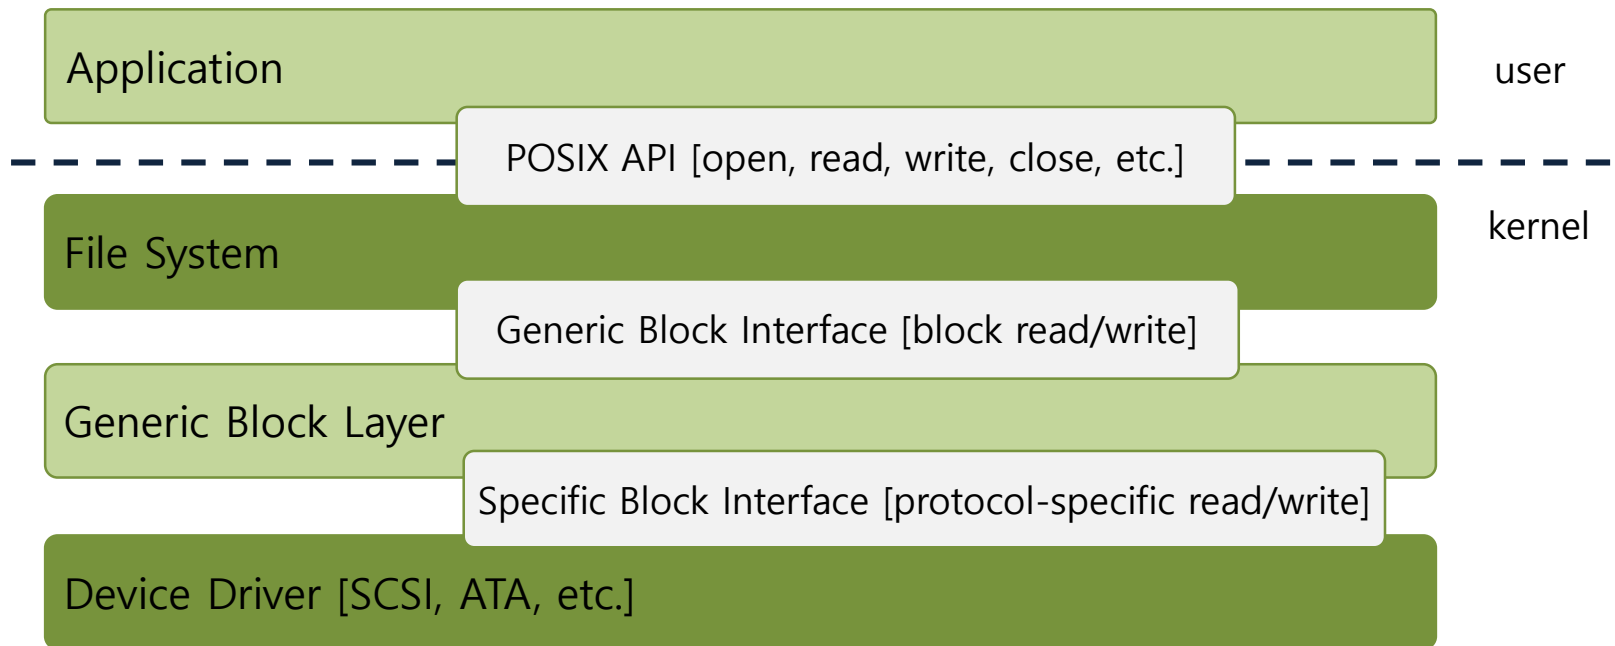
- How the OS communicates with the **device**?
- Solutions
 - **I/O instructions**: a way for the OS to send data to specific device registers.
 - Ex) `in` and `out` instructions on x86
 - **memory-mapped I/O**
 - Device registers available as if they were memory locations.
 - The OS `load` (to read) or `store` (to write) to the device instead of main memory.

Device interaction

- How the OS interact with devices **specific interfaces**?
 - Ex) We'd like to build a file system that worked on top of SCSI disks, IDE disks, USB keychain drivers, and so on.
- Solutions: **Abstraction**
 - Abstraction encapsulate **any specifics of device interaction**.

File system Abstraction

- File system **specifics** of which disk class it is using.
 - Ex) It issues **block read** and **write** request to the generic block layer.



The File System Stack

Problem of File System Abstraction

- If there is a device having many special capabilities, these capabilities **will go unused** in the generic interface layer.
- **Over 70% of OS** code is found in device drivers.
 - Many device drivers are needed because you might plug it to your system.
 - They are primary contributor to **kernel crashes**, making **more bugs**.

Reading Material

- **Chapter 36** of OSTEP book – by Remzi and Andrea Arpaci-Dusseau (University of Wisconsin)
<http://pages.cs.wisc.edu/~remzi/OSTEP/file-devices.pdf>

Questions?