

CIS 657 – Principles of Operating Systems

Topic: Concurrency – intro + thread API

Endadul Hoque

Acknowledgement

- Youjip Won (Hanyang University)
- OSTEP book – by Remzi and Andrea Arpaci-Dusseau (University of Wisconsin)

Thread

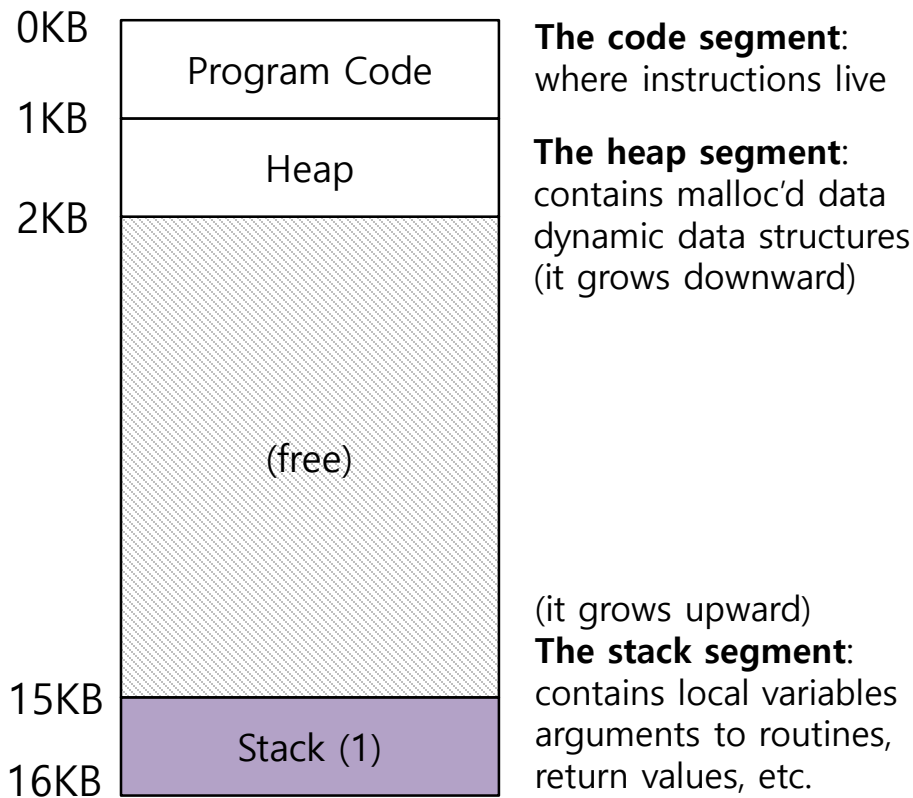
- A new abstraction for a single running process
- Multi-threaded program
 - A multi-threaded program has more than one point of execution.
 - Multiple PCs (Program Counters)
 - They **share** the same **address space**.

Context switch between threads


- Each thread has its **own** program counter and set of registers.
 - One or more **thread control blocks(TCBs)** are needed to store the state of each thread.
- When switching from running one (T1) to running the other (T2),
 - The register state of T1 will be saved.
 - The register state of T2 will be restored.
 - The **address space remains** the same.

Address Spaces

- There will be **one stack per thread** (**thread-local** storage)



**A Single-Threaded
Address Space**



**Two threaded
Address Space**

Why use Threads?

- Two major reasons
 - **Parallelism:**
 - A thread per CPU to speed up the process
 - **Avoid blocking** of progress due to slow I/O:
 - threading enables overlap of I/O with other activities within a single process
- Why not use multi-processes?
 - Threads **share an address space** and thus make it easy to share data
 - Using multi-processes is suitable for logically separate tasks, with little to no data sharing involved.

Example: Thread Traces

```
1  #include <stdio.h>
2  #include <assert.h>
3  #include <pthread.h>
4  #include "common.h"
5  #include "common_threads.h"
6
7  void *mythread(void *arg) {
8      printf("%s\n", (char *) arg);
9      return NULL;
10 }
11
12 int
13 main(int argc, char *argv[]) {
14     pthread_t p1, p2;
15     int rc;
16     printf("main: begin\n");
17     Pthread_create(&p1, NULL, mythread, "A");
18     Pthread_create(&p2, NULL, mythread, "B");
19     // join waits for the threads to finish
20     Pthread_join(p1, NULL);
21     Pthread_join(p2, NULL);
22     printf("main: end\n");
23     return 0;
24 }
```

Figure 26.2: Simple Thread Creation Code (t0.c)

Example: Thread Traces

main	Thread 1	Thread2
starts running prints "main: begin" creates Thread 1 creates Thread 2 waits for T1	runs prints "A" returns	
waits for T2		runs prints "B" returns
prints "main: end"		

Example: Thread Traces

	main	Thread 1	Thread2
m	starts running		
sta	prints "main: begin"		
pr	creates Thread 1		
cre		runs	
cre		prints "A"	
w		returns	
	creates Thread 2		
			runs
			prints "B"
w			returns
	waits for T1		
	<i>returns immediately; T1 is done</i>		
	waits for T2		
pr	<i>returns immediately; T2 is done</i>		
	prints "main: end"		

Example: Thread Traces

	main	Thread 1	Thread 2
m	starts running		
sta	prints "main: begin"		
pr	creates Thread 1		
cre	creates Thread 2		
cre			runs
wa			prints "B"
			returns
	waits for T1		
wa		runs	
		prints "A"	
		returns	
	waits for T2		
	<i>returns immediately; T2 is done</i>		
pr	prints "main: end"		

Sharing Data: Race Condition

- Example with two threads
 - **counter = counter + 1** (default is **50**)
 - We expect the result to be **52**. However,

```
100 mov 0x8049a1c, %eax
105 add $0x1, %eax
108 mov %eax, 0x8049a1c
```

OS	Thread1	Thread2	(after instruction) PC %eax counter
<hr/>			

Sharing Data: Race Condition

- Example with two threads
 - **counter = counter + 1** (default is **50**)
 - We expect the result to be **52**. However,

```

100 mov 0x8049a1c, %eax
105 add $0x1, %eax
108 mov %eax, 0x8049a1c
    
```

OS	Thread1	Thread2	(after instruction)		
			PC	%eax	counter
	<i>before critical section</i>		100	0	50
	mov 0x8049a1c, %eax		105	50	50
	add \$0x1, %eax		108	51	50
interrupt					
	<i>save T1's state</i>				
	<i>restore T2's state</i>		100	0	50
		mov 0x8049a1c, %eax	105	50	50
		add \$0x1, %eax	108	51	50
		mov %eax, 0x8049a1c	113	51	51
interrupt					
	<i>save T2's state</i>				
	<i>restore T1's state</i>		108	51	51
	mov %eax, 0x8049a1c		113	51	51

Sharing Data: Race Condition

- Example with two threads
 - **counter = counter + 1** (default is **50**)
 - We expect the result to be **52**. However,

```
100 mov 0x8049a1c, %eax
105 add $0x1, %eax
108 mov %eax, 0x8049a1c
```

(after instruction)

Race conditions (aka data race) can not only produce wrong results, but also produce different results across different runs – giving rise to indeterminate behavior

interrupt

save T2's state

restore T1's state

mov %eax, 0x8049a1c

108	51	50
113	51	51

Critical section

- A piece of code that **accesses a shared variable** (aka resources) and must **not be concurrently executed** by more than one thread.
 - Multiple threads executing critical section can result in a race condition.
 - Desired property: **mutual exclusion** guarantee for critical sections
 - One possible solution: **Atomic** execution of critical sections
 - Atomicity notion: "**all or none**"
- Based on a hardware synchronization primitives (i.e., special instructions), we can build complex abstractions to enable multi-threaded code produce correct results reliably.

Some interesting lines from Chapter 26

- It is a wonderful and hard problem, and should make your mind hurt (a bit).
- If it doesn't, then you don't understand!
- Keep working until your head hurts; you then **know** you're headed **in the right direction**.

Thread API

Thread Creation

- How to create and control threads?

```
#include <pthread.h>

int
pthread_create(      pthread_t*      thread,
                    const pthread_attr_t* attr,
                    void*             (*start_routine)(void*),
                    void*             arg);
```

- `thread`: Used to interact with this thread.
- `attr`: Used to specify any attributes this thread might have.
 - Stack size, Scheduling priority, ...
- `start_routine`: pointer to the function this thread will start executing from.
- `arg`: the argument to be passed to the function (`start_routine`)
 - a **void pointer** allows us to pass in **any type of** argument.

Thread Creation

```
#include <pthread.h>

typedef struct __myarg_t {
    int a;
    int b;
} myarg_t;

void *mythread(void *arg) {
    myarg_t *m = (myarg_t *) arg;
    printf("%d %d\n", m->a, m->b);
    return NULL;
}

int main(int argc, char *argv[]) {
    pthread_t p;
    int rc;

    myarg_t args;
    args.a = 10;
    args.b = 20;
    rc = pthread_create(&p, NULL, mythread, &args);
    ...
}
```

Wait for a thread to complete

```
int pthread_join(pthread_t thread, void **value_ptr);
```

- thread: Specify which thread *to wait for*
- value_ptr: A pointer to capture the return value
 - Because `pthread_join()` routine changes the value, you need to **pass in a pointer** to that value.

Wait for a thread to complete

```
1  #include <stdio.h>
2  #include <pthread.h>
3  #include <assert.h>
4  #include <stdlib.h>
5
6  typedef struct __myarg_t {
7      int a;
8      int b;
9  } myarg_t;
10
11 typedef struct __myret_t {
12     int x;
13     int y;
14 } myret_t;
15
16 void *mythread(void *arg) {
17     myarg_t *m = (myarg_t *) arg;
18     printf("%d %d\n", m->a, m->b);
19     myret_t *r = malloc(sizeof(myret_t));
20     r->x = 1;
21     r->y = 2;
22     return (void *) r;
23 }
24
```

Wait for a thread to complete

```
25  int main(int argc, char *argv[]) {
26      int rc;
27      pthread_t p;
28      myret_t *m;
29
30      myarg_t args;
31      args.a = 10;
32      args.b = 20;
33      pthread_create(&p, NULL, mythread, &args);
34      pthread_join(p, (void **) &m); // this thread has been
                                     // waiting inside of the
                                     // pthread_join() routine.
35      printf("returned %d %d\n", m->x, m->y);
36      return 0;
37  }

18  printf("%d %d\n", m->a, m->b);
19  myret_t *r = malloc(sizeof(myret_t));
20  r->x = 1;
21  r->y = 2;
22  return (void *) r;
23  }
24
```

Example: Dangerous code

- Be careful with how values are returned from a thread.

```
1  void *mythread(void *arg) {
2      myarg_t *m = (myarg_t *) arg;
3      printf("%d %d\n", m->a, m->b);
4      myret_t r; // ALLOCATED ON STACK: BAD!
5      r.x = 1;
6      r.y = 2;
7      return (void *) &r;
8  }
```

- When the variable `r` returns, it is automatically **de-allocated**.

Example: Simpler Argument Passing to a Thread

- Just passing in a single value

```
1  void *mythread(void *arg) {
2      int m = (int) arg;
3      printf("%d\n", m);
4      return (void *) (arg + 1);
5  }
6
7  int main(int argc, char *argv[]) {
8      pthread_t p;
9      int rc, m;
10     pthread_create(&p, NULL, mythread, (void *) 100);
11     pthread_join(p, (void **) &m);
12     printf("returned %d\n", m);
13     return 0;
14 }
```

Locks

- Provide **mutual exclusion** to a critical section

- Interface

```
int pthread_mutex_lock(pthread_mutex_t *mutex);  
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

- Usage (w/o *lock initialization* and *error check*)

```
pthread_mutex_t lock;  
pthread_mutex_lock(&lock);  
x = x + 1; // or whatever your critical section is  
pthread_mutex_unlock(&lock);
```

- No other thread holds the lock → the thread will acquire the lock and **enter the critical section**.
- If another thread hold the lock → the thread will **not return from the call** until it has acquired the lock.

Locks

- All locks must be properly initialized.
 - One way: using `PTHREAD_MUTEX_INITIALIZER`

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
```

- The dynamic way: using `pthread_mutex_init()`

```
int rc = pthread_mutex_init(&lock, NULL);  
assert(rc == 0); // always check success!
```

Locks

- Check errors code when calling lock and unlock
 - An example wrapper

```
// Use this to keep your code clean but check for failures
// Only use if exiting program is OK upon failure
void Pthread_mutex_lock(pthread_mutex_t *mutex) {
    int rc = pthread_mutex_lock(mutex);
    assert(rc == 0);
}
```

- These two calls are used in lock acquisition

```
int pthread_mutex_trylock(pthread_mutex_t *mutex);
int pthread_mutex_timelock(pthread_mutex_t *mutex,
                           struct timespec *abs_timeout);
```

- trylock: return failure if the lock is already held
- timedlock: return after a timeout

Condition Variables

- **Condition variables** are useful when some kind of **signaling** must take place between threads.

```
int pthread_cond_wait(pthread_cond_t *cond,  
                      pthread_mutex_t *mutex);  
int pthread_cond_signal(pthread_cond_t *cond);
```

– `pthread_cond_wait`:

- Put the calling thread to sleep.
- Wait for some other thread to signal it.

– `pthread_cond_signal`:

- Notify at least one of the threads that are blocked on the condition variable

Condition Variables

- A thread calling wait routine:

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;  
pthread_cond_t cond = PTHREAD_COND_INITIALIZER;  
  
pthread_mutex_lock(&lock);  
while (ready == 0)  
    pthread_cond_wait(&cond, &lock);  
pthread_mutex_unlock(&lock);
```

- The wait call **releases the lock** when putting said caller to sleep.
- Before returning after being woken, the wait call **re-acquire the lock**.

- A thread calling signal routine:

```
pthread_mutex_lock(&lock);  
ready = 1;  
pthread_cond_signal(&cond);  
pthread_mutex_unlock(&lock);
```

Condition Variables

- The waiting thread **re-checks** the condition **in a while loop**, instead of a simple if statement.

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;  
pthread_cond_t cond = PTHREAD_COND_INITIALIZER;  
  
pthread_mutex_lock(&lock);  
while (ready == 0)  
    pthread_cond_wait(&cond, &lock);  
pthread_mutex_unlock(&lock);
```

- Without rechecking, the waiting thread will continue thinking that the condition has changed *even though it has not*.

Condition Variables

- Poor way to do this.
 - A thread calling wait routine:

```
while (ready == 0)  
    ; // spin
```

- A thread calling signal routine:

```
ready = 1;
```

- It performs poorly in many cases. → just wastes CPU cycles.
 - It is error prone.

Compiling and Running

- To compile them, you must include the header `pthread.h`
 - Explicitly link with the **pthread library**, by adding the `-pthread` flag.

```
prompt> gcc -o main main.c -Wall -pthread
```

- For more information,

```
man -k pthread
```

Reading Material

- **Chapter 26-27** of OSTEP book – by Remzi and Andrea Arpaci-Dusseau (University of Wisconsin)
<http://pages.cs.wisc.edu/~remzi/OSTEP/threads-intro.pdf>
<http://pages.cs.wisc.edu/~remzi/OSTEP/threads-api.pdf>

Questions?