

CIS 657 – Principles of Operating Systems

Topic: Concurrency – Locks

Endadul Hoque

Acknowledgement

- Youjip Won (Hanyang University)
- OSTEP book – by Remzi and Andrea Arpaci-Dusseau (University of Wisconsin)

Locks: The basic idea

- Ensure that any **critical section** executes as if it were a **single atomic instruction**.
 - An example: the canonical update of a shared variable

```
balance = balance + 1;
```

- Add some code around the critical section

```
1  lock_t mutex; // some globally-allocated lock 'mutex'  
2  ...  
3  lock(&mutex);  
4  balance = balance + 1;  
5  unlock(&mutex);
```

Locks: The basic idea

- Lock variable holds the state of the lock.
 - **available** (or **unlocked** or **free**)
 - No thread holds the lock.
 - **acquired** (or **locked** or **held**)
 - Exactly one thread holds the lock and presumably is in a critical section.

The semantics of the lock()

- `lock()`
 - **Try to** acquire the lock.
 - If no other thread holds the lock, the thread will **acquire** the lock.
 - **Enter** the *critical section*.
 - This thread is said to be the owner of the lock.
 - Other threads are *prevented from* entering the critical section while the first thread that holds the lock is in there.

Pthread Locks - mutex

- The name that the POSIX library uses for a lock.
 - Used to provide **mutual exclusion** between threads.

```
1 pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;  
2  
3 Pthread_mutex_lock(&lock); // wrapper for pthread_mutex_lock()  
4 balance = balance + 1;  
5 Pthread_mutex_unlock(&lock);
```

- Instead of using one big lock (a **coarse-grained** approach), we may be using different locks to protect *different variables* (a more **fine-grained** approach) → Increase **concurrency**.

Building A Lock

- Efficient locks must provide mutual exclusion at **low cost**.
- Building a lock needs some help from the **hardware** and the **OS**.

Evaluating locks – Basic criteria

- **Correctness** (aka Mutual exclusion)
 - Does the lock work, preventing multiple threads from entering *a critical section*?
- **Fairness**
 - Does each thread contending for the lock get a fair shot at acquiring it once it is free? (Addressing starvation)
- **Performance**
 - The time overheads added by using the lock
 - **Cases:** (a) No contention, (b) contention on single CPU, (c) contention on multiple CPUs

Naïve solution: Controlling Interrupts

- **Disable Interrupts** for critical sections
 - One of the earliest solutions used to provide mutual exclusion
 - Invented for single-processor systems.

```
1  void lock() {  
2      DisableInterrupts();  
3  }  
4  void unlock() {  
5      EnableInterrupts();  
6  }
```

Naïve solution: Controlling Interrupts

- **Disable Interrupts** for critical sections
 - One of the earliest solutions used to provide mutual exclusion
 - Invented for single-processor systems.

```
1  void lock() {  
2      DisableInterrupts();  
3  }  
4  void unlock() {  
5      EnableInterrupts();  
6  }
```

- Problems:
 - Require too much *trust* in applications
 - Greedy (or malicious) program could monopolize the processor.
 - Do not work on **multiprocessors**
 - Disabling interrupts for a longer time can lose important interrupts
 - Code that masks or unmask interrupts be executed *slowly* by modern CPUs

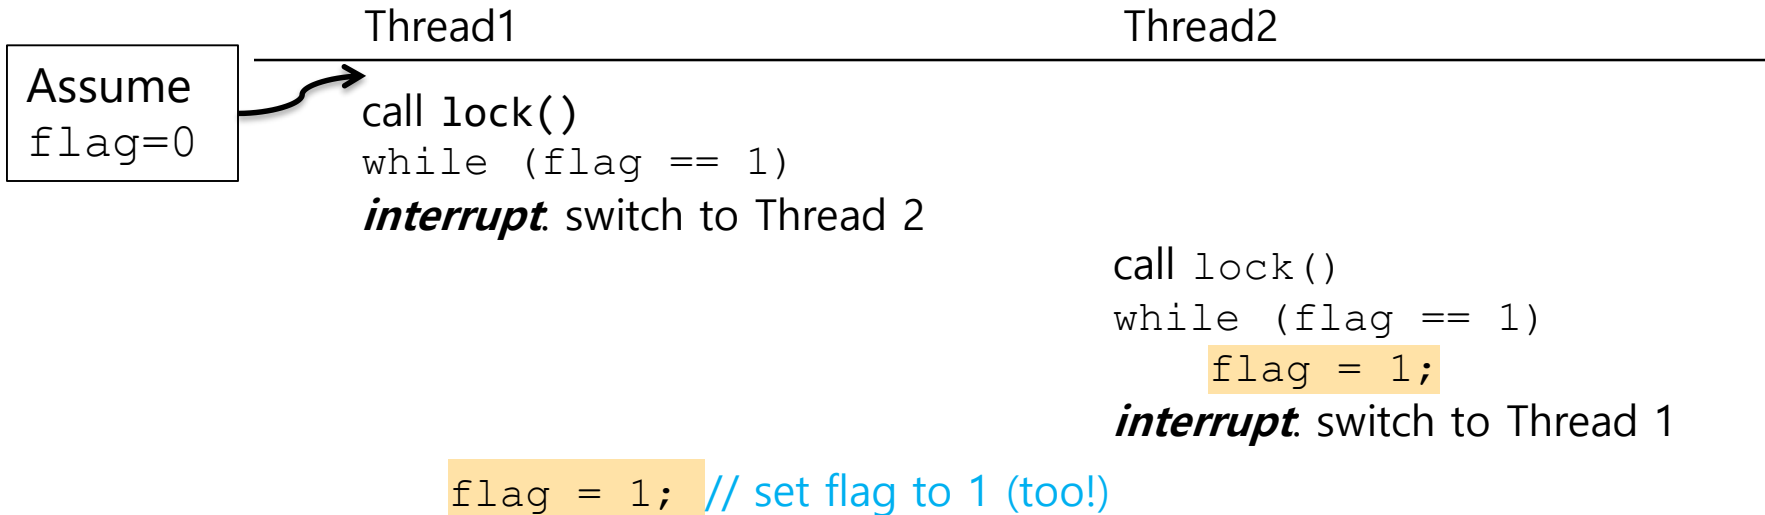
Why is hardware support needed?

- **First attempt:** Using a *flag* denoting whether the lock is held or not.
 - The code below has problems.

```
1  typedef struct __lock_t { int flag; } lock_t;
2
3  void init(lock_t *mutex) {
4      // 0 → lock is available, 1 → held
5      mutex->flag = 0;
6  }
7
8  void lock(lock_t *mutex) {
9      while (mutex->flag == 1) // TEST the flag
10         ; // spin-wait (do nothing)
11     mutex->flag = 1; // now SET it !
12 }
13
14 void unlock(lock_t *mutex) {
15     mutex->flag = 0;
16 }
```

Why is hardware support needed?

- **Problem 1:** No Mutual Exclusion



- **Problem 2:** Spin-waiting wastes time waiting for another thread.

- So, we need an atomic instruction supported by **Hardware!**
 - **E.g.,** *test-and-set* instruction, also known as atomic exchange

Test And Set (Atomic Exchange)

- An H/W instruction to support the creation of simple locks
- This instruction operates as follows (C pseudocode):

```
1  int TestAndSet(int *ptr, int new) {  
2      int old = *ptr; // fetch old value at ptr  
3      *ptr = new;     // store 'new' into ptr  
4      return old;     // return the old value  
5  }
```

- **return** old value pointed to by the `ptr` (testing)
- *Simultaneously* **update** said value to `new` (setting)
- This sequence of operations is **performed atomically**.

A Simple Spin Lock using test-and-set

```
1  typedef struct __lock_t {
2      int flag;
3  } lock_t;
4
5  void init(lock_t *lock) {
6      // 0 indicates that lock is available,
7      // 1 that it is held
8      lock->flag = 0;
9  }
10
11 void lock(lock_t *lock) {
12     while (TestAndSet(&lock->flag, 1) == 1)
13         ;    // spin-wait
14 }
15
16 void unlock(lock_t *lock) {
17     lock->flag = 0;
18 }
```

Note: To work correctly on *a single processor*, it requires a preemptive scheduler.

Evaluating Spin Locks (using Test-and-Set)

- **Correctness:** yes
 - The spin lock only allows a single thread to entry the critical section.
- **Fairness:** no
 - Spin locks don't provide any fairness guarantees.
 - Indeed, a thread spinning may spin *forever*.
- **Performance:**
 - In the single CPU, performance overheads can be quite *substantial*.
 - If the number of threads roughly equals the number of CPUs, spin locks work *reasonably well*.

Some other H/W instructions for Spin Locks

- Compare-and-Swap
- Load-Linked and Store-Conditional
- Fetch-And-Add

If interested to learn more on this, read sections 28.9, 28.10, 28.11 from the book chapter

So Much Spinning

- Hardware-based spin locks are **simple** and they work.
- In some cases, these solutions can be quite **inefficient**.
 - Any time a thread gets caught *spinning*, it **wastes an entire time slice** doing nothing but checking a value.

How To Avoid *Spinning*?
We'll need **OS Support** too!

A Simple Approach: Just Yield

- When you are going to spin, **give up the CPU** to another thread.
 - OS system call moves the caller from the *running state* to the *ready state*.
 - The cost of a **context switch** can be substantial and the **starvation** problem still exists.

```
1  void init() {
2      flag = 0;
3  }
4
5  void lock() {
6      while (TestAndSet(&flag, 1) == 1)
7          yield(); // give up the CPU
8  }
9
10 void unlock() {
11     flag = 0;
12 }
```

Lock with Test-and-set and Yield

Using Queues: Sleeping Instead of Spinning

- **Queue** to keep track of which threads are waiting to enter the lock.
- `park()`
 - Put a calling thread to sleep
- `unpark(threadID)`
 - Wake a particular thread as designated by `threadID`.

Using Queues: Sleeping Instead of Spinning

```
1  typedef struct __lock_t { int flag; int guard; queue_t *q; } lock_t;
2
3  void lock_init(lock_t *m) {
4      m->flag = 0;
5      m->guard = 0;
6      queue_init(m->q);
7  }
8
9  void lock(lock_t *m) {
10     while (TestAndSet(&m->guard, 1) == 1)
11         ; // acquire guard lock by spinning
12     if (m->flag == 0) {
13         m->flag = 1; // lock is acquired
14         m->guard = 0;
15     } else {
16         queue_add(m->q, gettid());
17         m->guard = 0;
18         park();
19     }
20 }
21 ...
```

Lock With Queues, Test-and-set, Yield, And Wakeup

Using Queues: Sleeping Instead of Spinning

```
22 void unlock(lock_t *m) {
23     while (TestAndSet(&m->guard, 1) == 1)
24         ; // acquire guard lock by spinning
25     if (queue_empty(m->q))
26         m->flag = 0; // let go of lock; no one wants it
27     else
28         unpark(queue_remove(m->q)); // hold lock (for next thread!)
29     m->guard = 0;
30 }
```

```
9 void lock(lock_t *m) {
10     while (TestAndSet(&m->guard, 1) == 1)
11         ; // acquire guard lock by spinning
12     if (m->flag == 0) {
13         m->flag = 1; // lock is acquired
14         m->guard = 0;
15     } else {
16         queue_add(m->q, gettid());
17         m->guard = 0;
18         park();
19     }
20 }
21 ...
```

Lock With Queues, Test-and-set, Yield, And Wakeup

Wakeup/waiting race

- Scenario:
 1. Thread A holds the lock
 2. Thread B attempts to acquire lock, hence ends up adding itself to the queue; but before it could call `park()`, context switch happens
 3. Thread A releases the lock and calls `unpark()`, but no thread to wake up
 4. Thread B regains the CPU and calls `park()` now; thus Thread B goes to sleep (and potentially sleeps forever).

Wakeup/waiting race

- **Solaris** solves this problem by adding a third system call: `setpark()`.
 - By calling this routine, a thread can indicate it *is about to* `park`.
 - If it happens to be interrupted and another thread calls `unpark` before `park` is actually called, the subsequent `park` returns immediately instead of sleeping.

```
1      queue_add(m->q, gettid());
2      setpark(); // new code
3      m->guard = 0;
4      park();
```

Code modification inside of `lock()`

Futex

- Linux provides a **futex** (is similar to Solaris's `park` and `unpark`), part of **nptl** library
 - `futex_wait(address, expected)`
 - Put the calling thread to sleep
 - If the value at `address` is not equal to `expected`, the call returns immediately.
 - `futex_wake(address)`
 - Wake one thread that is waiting on the queue.

Reading Material

- **Chapter 28** of OSTEP book – by Remzi and Andrea Arpaci-Dusseau (University of Wisconsin)
<http://pages.cs.wisc.edu/~remzi/OSTEP/threads-locks.pdf>

Questions?